

```
In [1]: 1 import pandas as pd
2 import numpy as np
3 from sklearn.model_selection import train_test_split
4 from sklearn.preprocessing import StandardScaler, LabelEncoder
5 from sklearn.ensemble import RandomForestClassifier
6 from sklearn.metrics import classification_report, confusion_matrix, ac
7 from imblearn.over_sampling import SMOTE
8 from sklearn.pipeline import Pipeline
```

```
In [2]: 1 column_names = ['duration', ' protocol_type',
2                  ' service', ' flag',
3                  ' src_bytes', ' dst_bytes',
4                  ' land', ' wrong_fragment',
5                  ' urgent', ' hot',
6                  ' num_failed_logins', ' logged_in',
7                  ' num_compromised', ' root_shell',
8                  ' su_attempted', ' num_root',
9                  ' num_file_creations', ' num_shells',
10                 ' num_access_files', ' num_outbound_cmds',
11                 ' is_host_login', ' is_guest_login',
12                 ' count', ' srv_count',
13                 ' serror_rate', ' srv_error_rate',
14                 ' rerror_rate', ' srv_rerror_rate',
15                 ' same_srv_rate', ' diff_srv_rate',
16                 ' srv_diff_host_rate', ' dst_host_count',
17                 ' dst_host_srv_count', ' dst_host_same_srv_rate',
18                 ' dst_host_diff_srv_rate', ' dst_host_same_src_port_rate',
19                 ' dst_host_srv_diff_host_rate', ' dst_host_serror_rate',
20                 ' dst_host_srv_rerror_rate', ' dst_host_rerror_rate',
21                 ' dst_host_srv_rerror_rate']
```

```
In [3]: 1 file_paths = [
2         'Data_of_Attack_Back.csv',
3         'Data_of_Attack_Back_BufferOverflow.csv',
4         'Data_of_Attack_Back_FTPWrite.csv',
5         'Data_of_Attack_Back_GuessPassword.csv',
6         'Data_of_Attack_Back_Neptune.csv',
7         'Data_of_Attack_Back_NMap.csv',
8         'Data_of_Attack_Back_Normal.csv',
9         'Data_of_Attack_Back_PortSweep.csv',
10        'Data_of_Attack_Back_RootKit.csv',
11        'Data_of_Attack_Back_Satan.csv',
12        'Data_of_Attack_Back_Smurf.csv',
13    ]
14
```

```
In [4]: 1
2 labels = ['Back', 'BufferOverflow', 'FTPWrite', 'GuessPassWord', 'Neptu
3
4
```

```
In [5]: 1 dataframes = []
2 for file_paths, label in zip(file_paths, labels):
3     if label == 'FTPWrite':
4         df = pd.read_csv(file_paths, header=None, names=column_names)
5     else:
6         df = pd.read_csv(file_paths)
7
8     df['Label'] = label
9     dataframes.append(df)
10
11
```


```
In [6]: 1 # Combine into a single DataFrame
2 df = pd.concat(dataframes, ignore_index=True)
```

```
In [7]: 1 df
```

Out[7]:

| | duration | protocol_type | service | flag | src_bytes | dst_bytes | land | wrong_fragment |
|--------|----------|---------------|---------|------|-----------|-----------|------|----------------|
| 0 | 0.0 | 0.00 | 0.00 | 0.0 | 0.54540 | 0.08314 | 0 | 0.0 |
| 1 | 0.0 | 0.00 | 0.00 | 0.0 | 0.54540 | 0.08314 | 0 | 0.0 |
| 2 | 0.0 | 0.00 | 0.00 | 0.0 | 0.54540 | 0.08314 | 0 | 0.0 |
| 3 | 0.0 | 0.00 | 0.00 | 0.0 | 0.54540 | 0.08314 | 0 | 0.0 |
| 4 | 0.0 | 0.00 | 0.00 | 0.0 | 0.54540 | 0.08314 | 0 | 0.0 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 817546 | 0.0 | 0.02 | 0.09 | 0.0 | 0.01032 | 0.00000 | 0 | 0.0 |
| 817547 | 0.0 | 0.02 | 0.09 | 0.0 | 0.01032 | 0.00000 | 0 | 0.0 |
| 817548 | 0.0 | 0.02 | 0.09 | 0.0 | 0.01032 | 0.00000 | 0 | 0.0 |
| 817549 | 0.0 | 0.02 | 0.09 | 0.0 | 0.01032 | 0.00000 | 0 | 0.0 |
| 817550 | 0.0 | 0.01 | 0.12 | 0.0 | 0.00028 | 0.00000 | 0 | 0.3 |

817551 rows × 42 columns



Data Preprocessing

```
In [8]: 1 # Drop rows with missing values or fill them
2 df.dropna(inplace=True) # or data.fillna(0, inplace=True)
3
4 # Split data into features and target
5 X = df.drop('Label', axis=1)
6 y = df['Label']
7
8 # Split data into training and testing sets
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
10
```

In [9]: 1 *#Handling Imbalanced Data*

In [10]: 1 smote = SMOTE()
2 X_train_sm, y_train_sm = smote.fit_resample(X_train, y_train)
3

In [11]: 1 *#Feature Scaling*
2 scaler = StandardScaler()
3 X_train_sm = scaler.fit_transform(X_train_sm)
4 X_test = scaler.transform(X_test)
5

In [12]: 1 *#Model Training*

In [13]: 1 clf = RandomForestClassifier()

In [14]: 1 clf.fit(X_train_sm, y_train_sm)

Out[14]: ▾ RandomForestClassifier
RandomForestClassifier()

Model Evaluation

In [15]: 1 y_pred = clf.predict(X_test)

```
In [16]: 1 # Print classification report and accuracy
2 print(classification_report(y_test, y_pred))
3 print("Accuracy:", accuracy_score(y_test, y_pred))
```

C:\Users\Asus\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\Asus\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

C:\Users\Asus\anaconda3\Lib\site-packages\sklearn\metrics_classification.py:1469: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

| | precision | recall | f1-score | support |
|----------------|-----------|--------|----------|---------|
| Back | 0.99 | 1.00 | 1.00 | 192 |
| BufferOverflow | 1.00 | 0.62 | 0.77 | 8 |
| FTPWrite | 1.00 | 1.00 | 1.00 | 1 |
| GuessPassWord | 1.00 | 1.00 | 1.00 | 9 |
| NMap | 0.99 | 1.00 | 1.00 | 301 |
| Neptune | 1.00 | 1.00 | 1.00 | 45575 |
| Normal | 1.00 | 1.00 | 1.00 | 115143 |
| PortSweep | 1.00 | 1.00 | 1.00 | 582 |
| RootKit | 0.00 | 0.00 | 0.00 | 3 |
| Satan | 1.00 | 1.00 | 1.00 | 1098 |
| Smurf | 1.00 | 1.00 | 1.00 | 599 |
| accuracy | | | 1.00 | 163511 |
| macro avg | 0.91 | 0.87 | 0.89 | 163511 |
| weighted avg | 1.00 | 1.00 | 1.00 | 163511 |

Accuracy: 0.9998899156631664

```
In [17]: 1 # Print confusion matrix
2 print("Confusion Matrix:")
3 print(confusion_matrix(y_test, y_pred))
```

Confusion Matrix:

```
[[ 192  0  0  0  0  0  0  0  0  0]
 [  0  5  0  0  0  0  3  0  0  0]
 [  0  0  1  0  0  0  0  0  0  0]
 [  0  0  0  9  0  0  0  0  0  0]
 [  0  0  0  0 301  0  0  0  0  0]
 [  0  0  0  0  0 45575  0  0  0  0]
 [  1  0  0  0  2  0 115137  1  0  2]
 [  0  0  0  0  0  0  1 581  0  0]
 [  0  0  0  0  0  0  2  0  0  1]
 [  0  0  0  0  0  0  4  1  0 1093]
 [  0  0  0  0  0  0  0  0  0  0]
 [599]]
```

Q1.Binomial classification: Detect anomalies by predicting Activity is normal or attack

```
In [18]: 1 #Create a Binary Target Variable
2 df['Binary_Label'] = df['Label'].apply(lambda x: 0 if x == 'Normal' else 1)
```

```
In [19]: 1 # Features (drop the 'Label' and 'Binary_Label' columns)
2 X = df.drop(['Label', 'Binary_Label'], axis=1)
```

```
In [20]: 1 # Binary target variable
2 y_bin = df['Binary_Label']
```

```
In [21]: 1 # Split data into training and testing sets
2 X_train_bin, X_test_bin, y_train_bin, y_test_bin = train_test_split(X,
```

```
In [22]: 1 #Train a Model
2 from sklearn.linear_model import LogisticRegression
```

```
In [23]: 1 # Initialize the Logistic Regression model
2 log_reg_bin = LogisticRegression(max_iter=1000) # Increase max_iter if
```

```
In [24]: 1 # Fit the model to the training data
        2 log_reg_bin.fit(X_train_bin, y_train_bin)
```

```
Out[24]: LogisticRegression
LogisticRegression(max_iter=1000)
```

```
In [25]: 1 # Predict on the testing set
        2 y_pred_bin = log_reg_bin.predict(X_test_bin)
```

```
In [26]: 1 #Evaluate the Model
        2
        3 # Print the classification report to see precision, recall, and F1-score
        4 print(classification_report(y_test_bin, y_pred_bin))
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.99 | 1.00 | 1.00 | 115143 |
| 1 | 1.00 | 0.99 | 0.99 | 48368 |
| accuracy | | | 0.99 | 163511 |
| macro avg | 1.00 | 0.99 | 0.99 | 163511 |
| weighted avg | 0.99 | 0.99 | 0.99 | 163511 |

```
In [27]: 1 # Print the confusion matrix
        2 print(confusion_matrix(y_test_bin, y_pred_bin))
```

```
[[114954   189]
 [   649 47719]]
```

Precision for both classes (normal activities and attacks) is very high, near or at 1.00, indicating that the model has a very high accuracy in predicting positive samples.

Recall is also impressive, especially for the normal activities (1.00), indicating that the model is almost perfect in identifying all the actual normal activities. For attacks, the recall is slightly lower (0.99), but still very high, indicating that the model identifies most of the actual attacks correctly.

F1-score, which is the harmonic mean of precision and recall, is near perfect for both classes, reinforcing the model's balanced performance in terms of precision and recall.

The confusion matrix further clarifies the results:

Out of 115,143 true normal activities, 114,954 were correctly classified as normal, with only 189 misclassified as attacks. Out of 48,368 true attacks, 47,719 were correctly identified, with 649 misclassified as normal activities. The accuracy of 0.99 suggests that the model correctly predicts the class for 99% of the cases in your test set.

```
In [ ]: 1
```


Q2 . Multinomial Classification: Detecting type of activity by predicting Activity is Normal or Back or Buffer Over flow or FTP Write or Guess Password or Neptune or N-Map or Port Sweep or Root Kit or Satan or Smurf

```
In [28]: 1 # Separate features and target variable
2 X = df.drop('Label', axis=1) # Drop the 'Label' column to get the feat
3 y = df['Label'] # Target variable is the activity type
```

```
In [29]: 1 # Split the data into training and testing sets
2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```
In [30]: 1 # Create a pipeline that first standardizes the data then applies the c
2 pipeline = Pipeline([
3     ('scaler', StandardScaler()),
4     ('classifier', RandomForestClassifier())
5 ])
```

```
In [31]: 1 # Train the model
2 pipeline.fit(X_train, y_train)
```



```
In [32]: 1 # Predict on the test set
2 y_pred = pipeline.predict(X_test)
```

```
1 # Evaluate the model
2 print("Classification Report:")
3 print(classification_report(y_test, y_pred))
4
5 print("Confusion Matrix:")
6 print(confusion_matrix(y_test, y_pred, labels=['Normal', 'Back', 'Buffer']))
7
```

| | precision | recall | f1-score | support |
|----------------|-----------|--------|----------|---------|
| Back | 1.00 | 1.00 | 1.00 | 192 |
| BufferOverflow | 1.00 | 0.88 | 0.93 | 8 |
| FTPWrite | 0.50 | 1.00 | 0.67 | 1 |
| GuessPassWord | 1.00 | 1.00 | 1.00 | 9 |
| NMap | 0.99 | 1.00 | 1.00 | 301 |
| Neptune | 1.00 | 1.00 | 1.00 | 45575 |
| Normal | 1.00 | 1.00 | 1.00 | 115143 |
| PortSweep | 1.00 | 1.00 | 1.00 | 582 |
| RootKit | 1.00 | 0.33 | 0.50 | 3 |
| Satan | 1.00 | 1.00 | 1.00 | 1098 |
| Smurf | 1.00 | 1.00 | 1.00 | 599 |

| | | | | |
|--------------|------|------|------|--------|
| accuracy | | | 1.00 | 163511 |
| macro avg | 0.95 | 0.93 | 0.92 | 163511 |
| weighted avg | 1.00 | 1.00 | 1.00 | 163511 |

| | | | | | | | | | |
|----------|-----|---|---|---|-------|---|---|---|------|
| [[115143 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 192 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 45575 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1096 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 599]] | | | | | | | | | |


```
In [34]: 1 # Print confusion matrix
          2 print("Confusion Matrix for Binary Classification:")
          3 print(confusion_matrix(y_test_bin, y_pred_bin))
          4
```

Confusion Matrix for Binary Classification:

```
[[114954  189]
 [   649 47719]]
```

```
In [35]: 1 # Print classification report
2 print("Classification Report for Multinomial Classification:")
3 print(classification_report(y_test, y_pred))
4
5 # Print confusion matrix
6 print("Confusion Matrix for Multinomial Classification:")
7 print(confusion_matrix(y_test, y_pred, labels=labels))
8
```

Classification Report for Multinomial Classification:

| | precision | recall | f1-score | support |
|----------------|-----------|--------|----------|---------|
| Back | 1.00 | 1.00 | 1.00 | 192 |
| BufferOverflow | 1.00 | 0.88 | 0.93 | 8 |
| FTPWrite | 0.50 | 1.00 | 0.67 | 1 |
| GuessPassword | 1.00 | 1.00 | 1.00 | 9 |
| NMap | 0.99 | 1.00 | 1.00 | 301 |
| Neptune | 1.00 | 1.00 | 1.00 | 45575 |
| Normal | 1.00 | 1.00 | 1.00 | 115143 |
| PortSweep | 1.00 | 1.00 | 1.00 | 582 |
| RootKit | 1.00 | 0.33 | 0.50 | 3 |
| Satan | 1.00 | 1.00 | 1.00 | 1098 |
| Smurf | 1.00 | 1.00 | 1.00 | 599 |
| accuracy | | | 1.00 | 163511 |
| macro avg | 0.95 | 0.93 | 0.92 | 163511 |
| weighted avg | 1.00 | 1.00 | 1.00 | 163511 |

Confusion Matrix for Multinomial Classification:

| | | | | | | | | | |
|--------|---|---|---|-------|-----|--------|-----|---|------|
| [[192 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 7 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 9 | 0 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 45575 | 0 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 301 | 0 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 115143 | 0 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 0 | 582 | 0 | 0 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 1 | 1 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 1 | 0 | 1 | 0 | 1096 |
| 0] | | | | | | | | | |
| [0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| 599]] | | | | | | | | | |

In []: 1

In []: 1

In []:

1

Overall Performance:

High Accuracy: The model achieves near-perfect accuracy (100%) across the board, which is excellent for a multiclass classification problem. **Precision and Recall:** For most activity types, both precision and recall are very high, often reaching 1.00, indicating the model's strong capability to correctly identify and classify different types of network activities.

Observations:

1.Neptune and Normal Activities: The model performs exceptionally well in identifying 'Neptune' and 'Normal' activities, which have the highest number of instances, with perfect precision and recall scores.

2.Buffer Overflow and RootKit: These categories have lower sample sizes and show some variation in recall scores ('BufferOverflow' at 0.88 and 'RootKit' at 0.33), suggesting the model may struggle slightly more with these less-represented classes.

3.FTPWrite: Despite having only one instance in the test set, the model identified it correctly, though the precision is lower (0.50) due to the model's overprediction in this category.

Areas for Improvement:

1.Handling Rare Classes: The variance in performance for 'BufferOverflow' and 'RootKit' points to potential challenges in handling rare classes. Techniques like oversampling, synthetic data generation (SMOTE), or cost-sensitive learning might improve performance in these categories.

2.FTPWrite Misclassification: The model's overprediction for 'FTPWrite' suggests a need for further investigation. It might be beneficial to explore feature relevance for this category or adjust class weighting to mitigate this bias.